

Unit and Fuzz Testing in Evaluating MongoDB Reliability

Youngjae Moon, Lisa Liu, Robert Sheng

Contents

- Project Overview
- Features
- Tools and Technologies
- VM Distribution
- VM Roles
- Test Metrics
- Team Member Roles
- Future Work
- Conclusion

Project Overview

- This project aims to simplify the deployment and evaluation of MongoDB by:
 - Automating the setup of MongoDB clusters using Docker and Kubernetes.
 - Providing a suite of unit tests to validate basic functionality of MongoDB interactions.
 - Integrating fuzz testing, a type of automated software testing, through Google Atheris to identify potential vulnerabilities.

Features

- Automated MongoDB Setup
 - Easily deploy a MongoDB cluster on Chameleon Cloud using Docker and Kubernetes.
- Unit Testing
 - Validate core MongoDB operations such as CRUD operations and replica set configurations.
- Fuzz Testing
 - Stress-test the MongoDB client or server with random or invalid inputs to identify crashes and vulnerabilities.
- Kubernetes StatefulSet
 - Deploy and manage MongoDB clusters on Kubernetes with persistent storage and replica sets.

Tools and Technologies

- Docker
 - For containerization of MongoDB.
- Kubernetes
 - For orchestrating MongoDB clusters.
- Python 3.12.6
 - For running benchmark and test scripts.
- Google Atheris
 - For fuzz testing.

VM Distribution

- This project employs Docker and Kubernetes across a four-VM setup to automate the deployment, configuration, & testing of MongoDB clusters.
- Each VM is designated for specific roles to maximize efficiency, scalability, and reliability.

VM Roles

- VM1: Control and Orchestration Node
 - Master node in Kubernetes.
 - Deploys all pods.
- VM3: Primary MongoDB Replica Set
 - Runs Kubernetes pod that hosts the primary MongoDB replica set.
- VM4: Secondary MongoDB Replica Set + MongoDB Client
 - Runs Kubernetes pod that hosts the secondary MongoDB replica set.
 - Runs Kubernetes pod that runs a Python MongoDB client to interact with the MongoDB pods.

VM Roles

- This VM distribution maximizes the project's flexibility by leveraging Docker for containerized MongoDB deployments, and Kubernetes for high-availability and clustered management.
- This setup facilitates comprehensive testing of MongoDB across different environments, ensuring reliability, scalability, and performance.

Test Metrics – Functional

- CRUD Operation Metrics
 - Insert Latency (ms): Measure the time taken to insert a single document or batch of documents.
 - Query Latency (ms): Evaluate the time taken to retrieve documents based on primary key lookup, range queries, compound key lookups, or text search queries (if indexes are enabled).
 - Update Latency (ms): Measure the time to update one or more fields in existing documents.
 - Delete Latency (ms): Measure the time to delete documents based on specific criteria.
- Replica Set Metrics
 - Replication Lag (ms): Time difference between when a write is applied to the primary node and when it is replicated to secondaries.

Test Metrics – Performance

- Throughput
 - Operations per Second (ops/s): Measure the number of CRUD operations MongoDB can process per second under a specific workload.
- Resource Utilisation
 - CPU Utilisation (%): Measure CPU usage under various workloads.
 - Memory Utilisation (%): Track memory usage during read and write-intensive operations.
 - Disk I/O (MB/s): Monitor disk reads/writes during heavy data insertions or queries.

Test Metrics – Reliability

- Durability
 - Data Loss on Failure: Measure the extent of data loss (if any) under sudden node or cluster failures.

Test Metrics – For Fuzz Testing

- Resilience
 - Crash Rate (% of operations): Percentage of operations causing crashes during fuzz testing.
- Vulnerability Metrics
 - Edge Case Coverage: Number and variety of edge cases detected during fuzz testing.
 - Execution Paths Tested: Percentage of code paths executed during fuzz testing.

Test Metrics – Benchmark

- Load Testing
 - Sustained Performance: Measure system performance over an extended period (e.g., 1-hour test).

Team Member Roles

- Robert Sheng
 - Writing codes in VM1, VM3 and VM4 for MongoDB setup through Docker and Kubernetes.
 - Generating various test metrics for evaluation.
- Lisa Liu
 - Writing codes in VM1, VM3 and VM4 for MongoDB setup through Docker and Kubernetes.
 - Generating various test metrics for evaluation.
- Youngjae Moon
 - Writing test codes in VM1
 - Making up slides and final report

Future Work

- CI/CD Integration
 - Automate deployment and testing pipeline using GitLab or GitHub Actions.
- Develop more sophisticated test metrics.
 - And achieve high coverage.
- Note: Demo will be done through recording a video.

Conclusion

- This project automates MongoDB deployment and testing, enhancing reliability and reducing setup time.
- By integrating unit and fuzz testing, it ensures MongoDB is robust under various conditions.

For more information...

[https://github.com/Pingumaniac/Unit_and_Fuzz_Testing_in_Evaluating_Mongo
DB_Reliability](https://github.com/Pingumaniac/Unit_and_Fuzz_Testing_in_Evaluating_Mongo_DB_Reliability)

Thank you!

Any questions?